

Divide-and-Conquer sequential Monte Carlo with applications to high dimensional filtering

Francesca R. Crucinio

11 May 2023

CSML Reading Group, Lancaster

- 1 Sequential Monte Carlo (SMC)
- 2 Divide-and-Conquer SMC (DaC-SMC)
- 3 High Dimensional Filtering

Sequential Monte Carlo (SMC)

Sequential Monte Carlo (SMC) methods approximate sequences of targets $\{\pi_t\}_{t \geq 1}$ of the form


$$\pi_t(x_{1:t}) \propto \pi_{t-1}(x_{1:t-1})M_t(x_{t-1}, x_t)G_t(x_t),$$

where M_t are Markov kernels and G_t are non-negative weight functions (Chopin and Papaspiliopoulos, 2020).

Can be used

- as an alternative to MCMC (Del Moral et al., 2006)
- for rare event simulation (Cérou et al., 2012)
- for ABC (Sisson et al., 2007)
- for filtering

State Space Models (SSM)



A state space model $(X_t, Y_t)_{t \geq 1} \subset \mathbb{R}^{d \times p}$ with

- transition density for the *latent* process $f_t(x_{t-1}, x_t)$
- likelihood for the *observation* process $g_t(x_t, y_t)$

Consider


$$\begin{aligned} p(x_{1:t} | y_{1:t}) &\propto \prod_{k=1}^t f_k(x_{k-1}, x_k) g_k(x_k, y_k) \\ &= p(x_{1:t-1}, y_{1:t-1}) f_t(x_{t-1}, x_t) g_t(x_t, y_t). \end{aligned}$$

Sequential Monte Carlo (SMC)

- **Aim:** Approximate $\pi_t(x_t)$ using a set of N weighted particles $\{X_t^i, W_t^i\}_{i=1}^N$
- **Ingredients:**
 1. a Markov kernel to sample $X_t^i \sim M_t(\tilde{X}_{t-1}^i, \cdot)$
 2. a positive weight function to compute the weights $W_t^i \propto G_t(X_t^i)$
 3. a resampling scheme to obtain $\{\tilde{X}_t^i, 1/N\}_{i=1}^N$ from $\{X_t^i, W_t^i\}_{i=1}^N$

Example: the bootstrap particle filter corresponds to $\pi_t(x_t) = p(x_t|y_{1:t})$, $M_t = f_t$ and $G_t = g(y_t|\cdot)$.


Sequential Monte Carlo (SMC)



- 1 Sequential Monte Carlo (SMC)
- 2 Divide-and-Conquer SMC (DaC-SMC)**
- 3 High Dimensional Filtering

Divide-and-Conquer SMC (DaC-SMC)

An extension of standard SMC in which the sequence of target distributions $\{\pi_u\}_{u \in \mathbb{T}}$ evolves on a tree \mathbb{T} rather than on a line (Lindsten et al., 2017).



For simplicity we will also consider the unnormalised targets $\{\gamma_u\}_{u \in \mathbb{T}}$, s.t.
 $\pi_u = \gamma_u / Z_u$.


Divide-and-Conquer SMC (DaC-SMC)

- easier to distribute/parallelise (Corneflores et al., 2022; Ding and Gandy, 2018)
- some models more naturally adapted to a tree structure than a ‘linear’ one: e.g. graphical models (Lindsten et al., 2017; Paige and Wood, 2016; Jewell, 2015)

Level 0:

Level 1:

Level 2:



- distributed inference (Chan et al., 2021; Manderson and Goudie, 2021)

Divide-and-Conquer SMC (DaC-SMC)

At each node u we have a set of N weighted particles $\{X_u^i, W_u^i\}_{i=1}^N$ approximating π_u .

To evolve use standard SMC ingredients

- transition/mutation: $X_u^i \sim M_u((X_{\ell(u)}^i, X_{r(u)}), \cdot)$
- reweighting: $W_u^i \propto G_u(X_u^i)$
- resampling

with the addition of a **merging** step when two branches of the tree merge.

The Merge Step

- **Aim:** Given two populations of weighted particles on the left and the right child, $\{X_{\ell(u)}^i, W_{\ell(u)}^i\}_{i=1}^N$ and $\{X_{r(u)}^i, W_{r(u)}^i\}_{i=1}^N$, approximating $\gamma_{\ell(u)}$ and $\gamma_{r(u)}$ build an approximation of γ_u

The Merge Step

- **Aim:** Given two populations of weighted particles on the left and the right child, $\{X_{\ell(u)}^i, W_{\ell(u)}^i\}_{i=1}^N$ and $\{X_{r(u)}^i, W_{r(u)}^i\}_{i=1}^N$, approximating $\gamma_{\ell(u)}$ and $\gamma_{r(u)}$ build an approximation of γ_u

- **Ingredients:**

- ▶ the (weighted) product form estimator ([Kuntz et al., 2022](#))

$$\gamma_{\ell(u)}^N \times \gamma_{r(u)}^N = \frac{1}{N^2} \sum_{i_1=1}^N \sum_{i_2=1}^N W_{\ell(u)}^{i_1} W_{r(u)}^{i_2} \delta_{(X_{\ell(u)}^{i_1}, X_{r(u)}^{i_2})}$$

- ▶ N^2 mixture (importance) weights

$$m_u^{(i_1, i_2)} := W_{\ell(u)}^{i_1} W_{r(u)}^{i_2} \frac{\gamma_u(X_{\ell(u)}^{i_1}, X_{r(u)}^{i_2})}{\gamma_{\ell(u)}(X_{\ell(u)}^{i_1}) \gamma_{r(u)}(X_{r(u)}^{i_2})}$$

- ▶ a resampling scheme to obtain $\{\tilde{X}_u^i, 1/N\}_{i=1}^N$ from $\{(X_{\ell(u)}^{i_1}, X_{r(u)}^{i_2}), m_u^{(i_1, i_2)}\}_{i_1, i_2=1}^N$

Cheaper Alternatives

- lightweight mixture resampling ([Lindsten et al., 2017](#))
- lazy resampling schemes ([Corneflos et al., 2022](#))
- strategies borrowed from the literature on incomplete U-statistics ([Kuntz et al., 2021](#))

Lightweight Mixture Resampling (Lindsten et al., 2017)

- Fix $\theta \ll N$.
- Sample $\{X_{\ell(u)}^{i_1}\}_{i_1=1}^{\theta}$ from $\gamma_{\ell(u)}$ and $\{X_{r(u)}^{i_2}\}_{i_2=1}^{\theta}$ from $\gamma_{r(u)}$.
- Build the θ pairs $(X_{\ell(u)}^{i_1}, X_{r(u)}^{i_2})$ and compute their weights $m_u^{(i_1, i_2)}$.

Adaptive Lightweight Mixture Resampling

A strategy based on the effective sample size $(\sum_n m_u^n)^2 / \sum_n (m_u^n)^2$.

- ▶ Set a target ESS, ESS^* .
- ▶ For $i = 1, \dots, N$ build $(X_{\ell(u)}^i, X_{r(u)}^i)$ and compute their weights $m_u^{(i,i)}$.
- ▶ While $\text{ESS} < \text{ESS}^*$:
 - Draw one permutation of N , $\pi(N)$, build $(X_{\ell(u)}^i, X_{r(u)}^{\pi(i)})$ and compute their weights $m_u^{(i,\pi(i))}$.
 - Update ESS.

Adaptive Lightweight Mixture Resampling




Figure 1: Distribution of the number of permutations selected for a simple linear Gaussian model with $d = 32$, $N = 10^3$, and 10 time steps; the panels correspond to the levels of the tree from the level above the leaves (level 1) to the root (level 5).

Theoretical Properties (Kuntz et al., 2021)

- strong law of large numbers

$$\frac{1}{N} \sum_{i=1}^N \varphi(\tilde{X}_u^i) \xrightarrow{a.s.} \int \varphi(x) \pi_u(x) dx$$

- central limit theorem


$$\sqrt{N} \left(\frac{1}{N} \sum_{i=1}^N \varphi(\tilde{X}_u^i) - \int \varphi(x) \pi_u(x) dx \right) \xrightarrow{d} \mathcal{N}(0, \sigma_u^2(\varphi))$$

- unbiasedness of normalising constant estimates

And other: \mathbb{L}^p errors, bias estimates, optimal intermediate targets, optimal proposals, ...

- 1 Sequential Monte Carlo (SMC)
- 2 Divide-and-Conquer SMC (DaC-SMC)
- 3 High Dimensional Filtering**
 - Main Ideas
 - Experiments

State Space Models (SSM)




A state space model $(X_t, Y_t)_{t \geq 1} \subset \mathbb{R}^{d \times p}$ with

- transition density for the *latent* process $f_t(x_{t-1}, x_t)$
- likelihood for the *observation* process $g_t(x_t, y_t)$

We are interested in

$$\begin{aligned} p(x_{1:t} | y_{1:t}) &\propto \prod_{k=1}^t f_k(x_{k-1}, x_k) g_k(x_k, y_k) \\ &= p(x_{1:t-1}, y_{1:t-1}) f_t(x_{t-1}, x_t) g_t(x_t, y_t). \end{aligned}$$

Spatial SSM



1 Sequential Monte Carlo (SMC)

2 Divide-and-Conquer SMC (DaC-SMC)

3 High Dimensional Filtering

- Main Ideas

- Experiments

Divide and Conquer SMC for Filtering




Figure 2: Space decomposition for $d = 8$.

Choice of Distributions

Fix t . We need $\{\gamma_{t,u}\}_{u \in \mathbb{T}}$.

Constraints:

- $\gamma_{t,\mathfrak{R}}(x_t) \propto p(x_t|y_{1:t})$
- for $u \neq \mathfrak{R}$, $\gamma_{t,u}$ only depends on a subset of $x_t(1:d)$
- we can compute

$$\frac{\gamma_{t,u}(x_{t,\ell(u)}, x_{t,r(u)})}{\gamma_{t,\ell(u)}(x_{t,\ell(u)})\gamma_{t,r(u)}(x_{t,r(u)})}$$

Choice of Distributions

Recall

$$p(x_t|y_{1:t}) \propto g_t(x_t, y_t) \int f_t(x_{t-1}, x_t) p(x_{t-1}|y_{1:t-1}) dx_{t-1}.$$

Use auxiliary functions $f_{t,u}, g_{t,u}$ such that $f_{t,\mathfrak{R}} = f_t, g_{t,\mathfrak{R}} = g_t$ and for $u \in \mathbb{T} \setminus \mathfrak{R}$, $f_{t,u}$ and $g_{t,u}$ serve as proxies for marginals of f_t, g_t

$$\gamma_{t,u}(x_{t,u}) = g_{t,u}(x_{t,u}, (y_t(i))_{i \in u}) \int f_{t,u}(x_{t-1}, x_{t,u}) \gamma_{t-1,\mathfrak{R}}(x_{t-1}) dx_{t-1}$$

Choice of Distributions

Recall

$$p(x_t|y_{1:t}) \propto g_t(x_t, y_t) \int f_t(x_{t-1}, x_t) p(x_{t-1}|y_{1:t-1}) dx_{t-1}.$$

Use auxiliary functions $f_{t,u}, g_{t,u}$ such that $f_{t,\mathfrak{R}} = f_t, g_{t,\mathfrak{R}} = g_t$ and for $u \in \mathbb{T} \setminus \mathfrak{R}$, $f_{t,u}$ and $g_{t,u}$ serve as proxies for marginals of f_t, g_t

$$\gamma_{t,u}(x_{t,u}) = g_{t,u}(x_{t,u}, (y_t(i))_{i \in u}) \int f_{t,u}(x_{t-1}, x_{t,u}) \gamma_{t-1,\mathfrak{R}}(x_{t-1}) dx_{t-1}$$

The integral above is intractable, so we define approximate targets as in marginal particle filters ([Klaas et al., 2005](#))

$$\gamma_{t,u}(x_{t,u}) = g_{t,u}(x_{t,u}, (y_t(i))_{i \in u}) W_{t-1}^i \sum_{i=1}^N f_{t,u}(X_{t-1,\mathfrak{R}}^i, x_{t,u})$$

Divide and Conquer SMC for Filtering

At time t :

- start with $\{X_{t-1,\mathfrak{R}}^i, W_{t-1}^i\}_{i=1}^N$ approximating $p(x_{t-1}|y_{1:t-1})$ at the leaf level
- define targets

$$\gamma_{t,u}(x_{t,u}) = g_{t,u}(x_{t,u}, (y_t(i))_{i \in u}) W_{t-1}^i \sum_{i=1}^N f_{t,u}(X_{t-1,\mathfrak{R}}^i, x_{t,u})$$

- use DaC to move up the space from the leaves to the root, i.e. from d particle populations approximating 1-dimensional to one d -dimensional population $\{X_{t,\mathfrak{R}}^i, W_t^i\}_{i=1}^N$ which approximates $p(x_t|y_{1:t})$

Pros and Cons

Pros

- No need for analytical form of $p(x_t(i) \mid x_t(1 : i - 1))$
- No need for factorised likelihoods
- Easy to parallelise and distribute

Cons

- Polynomial cost in N (can be mitigated via GPUs)
- Needs specification of $\gamma_{t,u}$

1 Sequential Monte Carlo (SMC)

2 Divide-and-Conquer SMC (DaC-SMC)

3 High Dimensional Filtering

- Main Ideas

- Experiments


$$\begin{aligned}f_1(x_1) &= \mathcal{N}(x_1; m_1, \Sigma_1) \\f_t(x_{t-1}, x_t) &= \mathcal{N}(x_t; Ax_{t-1}, \Sigma) \\g_t(x_t, y_t) &= \mathcal{N}(y_t; x_t, \sigma_y^2 Id_d),\end{aligned}$$

with $\Sigma \in \mathbb{R}^{d \times d}$ a tridiagonal matrix.

We compare the RMSE

$$\text{RMSE}(x_t(i)) := \frac{\mathbb{E} [(\bar{x}_t(i) - \mu_{t,i})^2]}{\sigma_{t,i}^2}$$

Linear Gaussian Model - RMSE



Spatial Model

The components of X_t are indexed by the vertices $v \in V$ of a lattice, where $V = \{1, \dots, d\}^2$, and


$$X_t(v) = X_{t-1}(v) + U_t(v), \quad U_t(v) \sim \mathcal{N}(0, \sigma_x^2)$$

and

$$g_t(x_t, y_t) \propto$$


$$\left[1 + \nu^{-1} \sum_{v \in V} \left((y_t(v) - x_t(v)) \sum_{j: D(v,j) \leq r_y} \tau^{D(v,j)} (y_t(j) - x_t(j)) \right) \right]^{-(\nu + d^2)/2}.$$

2×2 grid




8×8 grid

(1, 1)



Runtime / s

(8, 6)



Runtime / s

 dac  dac-ada

Multivariate Stochastic Volatility Model

$$\begin{aligned}X_1 &\sim \mathcal{N}_d(0, \Sigma_0), \\Y_t &= C_t^{1/2} V_t, \\X_{t+1} &= \Phi X_t + U_t,\end{aligned}$$

with C_t a diagonal matrix with entries


$$C_t = \text{diag}(\exp(X_t(1)), \exp(X_t(2)), \dots, \exp(X_t(d))),$$

Φ a diagonal matrix with diagonal entries ϕ_i ,

$$\begin{pmatrix} V_t \\ U_t \end{pmatrix} \sim \mathcal{N}_{2d}(0, \Sigma) \quad \text{with } \Sigma = \begin{pmatrix} \Sigma_{VV} & \Sigma_{VU} \\ \Sigma_{VU} & \Sigma_{UU} \end{pmatrix},$$

and Σ_0 has entries given by

$$(\Sigma_0)_{ij} = (\Sigma_{UU})_{ij} / (1 - \phi_i \phi_j),$$



Opportunities and Related Ideas

- Theoretical guarantees should follow from those of standard DaC-SMC (Kuntz et al., 2021) and marginal particle filters (Crucinio and Johansen, 2023)
- smoothing, parameter estimation
- Other applications: Parallel (in time) Smoothing (Corneflos et al., 2022; Ding and Gandy, 2018), Divide and Conquer Fusion (Chan et al., 2021), inference for graphical models and phylogenetic trees (Jewell, 2015; Paige and Wood, 2016; Lindsten et al., 2017)

Thank you!

Bibliography i

- Frédéric Cérou, Pierre Del Moral, Teddy Furon, and Arnaud Guyader. Sequential monte carlo for rare event estimation. *Statistics and computing*, 22(3):795–808, 2012.
- R. Chan, M. Pollock, A. M. Johansen, and G. O. Roberts. Divide-and-conquer Monte Carlo fusion. *arXiv preprint arXiv:2110.07265*, 2021.
- N Chopin and O Papaspiliopoulos. *An Introduction to Sequential Monte Carlo*. Springer, Cham, 2020. doi: 10.1007/978-3-030-47845-2.
- A. Corneflos, N. Chopin, and S. Särkkä. De-Sequentialized Monte Carlo: a parallel-in-time particle smoother. *Journal of Machine Learning Research*, 23(283):1–39, 2022.
- F. R. Crucinio and A. M. Johansen. Properties of marginal sequential Monte Carlo methods. *arXiv preprint arXiv:2303.03498*, 2023.
- Pierre Del Moral, Arnaud Doucet, and Ajay Jasra. Sequential Monte Carlo samplers. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 68(3):411–436, 2006.
- D. Ding and A. Gandy. Tree-based particle smoothing algorithms in a hidden markov model. *arXiv preprint arXiv:1808.08400*, 2018.
- S. W. Jewell. *Divide and conquer sequential Monte Carlo for phylogenetics*. Master thesis, University of British Columbia, 2015.
- Mike Klaas, Nando De Freitas, and Arnaud Doucet. Toward practical N^2 Monte Carlo: The marginal particle filter. In *Proceedings of the 21st Conference in Uncertainty in Artificial Intelligence*, pages 308–315, 2005.
- J. Kuntz, F. R. Crucinio, and A. M. Johansen. The divide-and-conquer sequential Monte Carlo algorithm: theoretical properties and limit theorems. *arXiv preprint arXiv:2110.15782*, 2021.
- Juan Kuntz, Francesca R Crucinio, and Adam M Johansen. Product-form estimators: exploiting independence to scale up Monte Carlo. *Statistics and Computing*, 32(1):1–22, 2022.
- Fredrik Lindsten, Adam M Johansen, Christian A Næseth, Bonnie Kirkpatrick, Thomas B Schön, John A D Aston, and Alexandre Bouchard-Côté. Divide-and-Conquer with sequential Monte Carlo. *Journal of Computational and Graphical Statistics*, 26(2):445–458, 2017.
- Andrew A Manderson and Robert JB Goudie. Combining chains of bayesian models with markov melding. *arXiv preprint arXiv:2111.11566*, 2021.
- B. Paige and F. Wood. Inference networks for sequential monte carlo in graphical models. In *33rd Int Conf Mach Learn*, volume 48, pages 3040–3049, 2016.
- S. A. Sisson, Y. Fan, and M. M. Tanaka. Sequential Monte Carlo without likelihoods. 104(4):1760–1765, February 2007.