

Divide-and-Conquer sequential Monte Carlo with applications to high dimensional filtering

Francesca R. Crucinio, King's College London

ISBA Satellite workshop

Joint work with Adam M. Johansen

- 1 Motivation
- 2 Sequential Monte Carlo and Particle Filters
- 3 Divide and Conquer SMC
- 4 Divide and Conquer SMC for Filtering

Spatio-temporal Data





Figure 1: Drought Detection from Rain Precipitations (Næsseth et al., 2015).

- **Observe** precipitation (in millimeters) for each location and year
- **Recover** if there is a drought in that location

State Space Models (SSM)




A state space model $(X_n, Y_n)_{n \geq 1} \subset \mathbb{R}^{d \times p}$ with

- transition density for the *latent* process $f_n(x_{n-1}, x_n)$
- likelihood for the *observation* process $g_n(x_n, y_n)$

We are interested in


$$p(x_{1:n} | y_{1:n}) \propto p(x_{1:n-1}, y_{1:n-1}) f_n(x_{n-1}, x_n) g_n(x_n, y_n).$$

Spatial State Space Model



- 1 Motivation
- 2 Sequential Monte Carlo and Particle Filters**
- 3 Divide and Conquer SMC
- 4 Divide and Conquer SMC for Filtering

Sequential Monte Carlo (SMC)



Sequential Monte Carlo (SMC)


- **Aim:** Approximate $p(x_n|y_{1:n})$ using a set of N weighted particles $\{X_n^i, W_n^i\}_{i=1}^N$
- **Ingredients:**
 1. a transition density to sample $X_n^i \sim f_n(\tilde{X}_{n-1}^i, \cdot)$
 2. a likelihood function to compute the weights $W_n^i \propto g_n(X_n^i)$
 3. a resampling scheme to obtain $\{\tilde{X}_n^i, 1/N\}_{i=1}^N$ from $\{X_n^i, W_n^i\}_{i=1}^N$

Also known as **bootstrap particle filter**.

- 1 Motivation
- 2 Sequential Monte Carlo and Particle Filters
- 3 Divide and Conquer SMC**
- 4 Divide and Conquer SMC for Filtering

Divide-and-Conquer SMC (DaC-SMC)

An extension of standard SMC in which the sequence of target distributions $\{\hat{\gamma}_u\}_{u \in \mathbb{T}}$ evolves on a tree \mathbb{T} rather than on a line (Lindsten et al., 2017).



Divide-and-Conquer SMC (DaC-SMC)

At each node u we have a set of N weighted particles $\{X_u^i, W_u^i\}_{i=1}^N$ approximating $\hat{\gamma}_u$.

To evolve use standard SMC ingredients

- transition/mutation: $X_u^i \sim M_u((X_{\ell(u)}^i, X_{r(u)}^i), \cdot)$
- reweighting: $W_u^i \propto G_u(X_u^i)$
- resampling

with the addition of a **merging** step when two branches of the tree merge.

The Merge Step

- **Aim:** Given two populations of weighted particles on the left and the right child, $\{X_{\ell(u)}^i, W_{\ell(u)}^i\}_{i=1}^N$ and $\{X_{r(u)}^i, W_{r(u)}^i\}_{i=1}^N$, approximating $\hat{\gamma}_{\ell(u)}$ and $\hat{\gamma}_{r(u)}$ build an approximation of $\hat{\gamma}_u$

The Merge Step

- **Aim:** Given two populations of weighted particles on the left and the right child, $\{X_{\ell(u)}^i, W_{\ell(u)}^i\}_{i=1}^N$ and $\{X_{r(u)}^i, W_{r(u)}^i\}_{i=1}^N$, approximating $\hat{\gamma}_{\ell(u)}$ and $\hat{\gamma}_{r(u)}$ build an approximation of $\hat{\gamma}_u$

- **Ingredients:**

- ▶ the (weighted) product form estimator ([Kuntz et al., 2022](#))

$$\hat{\gamma}_{\ell(u)}^N \times \hat{\gamma}_{r(u)}^N = \frac{1}{N^2} \sum_{i_1=1}^N \sum_{i_2=1}^N W_{\ell(u)}^{i_1} W_{r(u)}^{i_2} \delta_{(X_{\ell(u)}^{i_1}, X_{r(u)}^{i_2})}$$

- ▶ N^2 mixture (importance) weights

$$m_u^{(i_1, i_2)} := W_{\ell(u)}^{i_1} W_{r(u)}^{i_2} \frac{\hat{\gamma}_u(X_{\ell(u)}^{i_1}, X_{r(u)}^{i_2})}{\hat{\gamma}_{\ell(u)}(X_{\ell(u)}^{i_1}) \hat{\gamma}_{r(u)}(X_{r(u)}^{i_2})}$$

- ▶ a resampling scheme to obtain $\{\tilde{X}_u^i, 1/N\}_{i=1}^N$ from $\{(X_{\ell(u)}^{i_1}, X_{r(u)}^{i_2}), m_u^{(i_1, i_2)}\}_{i_1, i_2=1}^N$

- 1 Motivation
- 2 Sequential Monte Carlo and Particle Filters
- 3 Divide and Conquer SMC
- 4 Divide and Conquer SMC for Filtering**

Divide and Conquer SMC for Filtering




Figure 2: Space decomposition for $d = 8$.

Divide and Conquer SMC for Filtering

At time n :

- start with $\{X_{n-1,\mathfrak{R}}^i, W_{n-1}^i\}_{i=1}^N$ approximating $p(x_{n-1}|y_{1:n-1})$ at the leaf level
- define targets

$$\hat{\gamma}_{n,u}(x_{n,u}) = g_{n,u}(x_{n,u}, (y_n(i))_{i \in u}) W_{n-1}^i \sum_{i=1}^N f_{n,u}(X_{n-1,\mathfrak{R}}^i, x_{n,u})$$

- use DaC to move up the space from the leaves to the root, i.e. from d particle populations approximating 1-dimensional to one d -dimensional population $\{X_{n,\mathfrak{R}}^i, W_n^i\}_{i=1}^N$ which approximates $p(x_n|y_{1:n})$

Linear Gaussian Model


$$\begin{aligned}f_1(x_1) &= \mathcal{N}(x_1; m_1, \Sigma_1) \\f_n(x_{n-1}, x_n) &= \mathcal{N}(x_n; Ax_{n-1}, \Sigma) \\g_n(x_n, y_n) &= \mathcal{N}(y_n; x_n, \sigma_y^2 Id_d),\end{aligned}$$

with $\Sigma \in \mathbb{R}^{d \times d}$ a tridiagonal matrix.

We compare the RMSE

$$\text{RMSE}(x_n(i)) := \frac{\mathbb{E}[(\bar{x}_n(i) - \mu_{n,i})^2]}{\sigma_{n,i}^2}$$

Linear Gaussian Model - RMSE



Spatial Model

The components of X_t are indexed by the vertices $v \in V$ of a lattice, where $V = \{1, \dots, d\}^2$, and


$$X_n(v) = X_{n-1}(v) + U_n(v), \quad U_n(v) \sim \mathcal{N}(0, \sigma_x^2)$$

and

$$g_t(x_n, y_n) \propto$$


$$\left[1 + \nu^{-1} \sum_{v \in V} \left((y_n(v) - x_n(v)) \sum_{j: D(v,j) \leq r_y} \tau^{D(v,j)} (y_n(j) - x_n(j)) \right) \right]^{-(\nu + d^2)/2}.$$

2×2 grid




8×8 grid

(1, 1)



Runtime / s

(8, 6)



Runtime / s

 dac  dac-ada

Pros

- No need for analytical form of $f_n(x_n(i) \mid x_n(1 : i - 1))$
- No need for factorised likelihoods
- Easy to parallelise and distribute

Cons

- Polynomial cost in N (can be mitigated via GPUs)
- Needs specification of $\hat{\gamma}_{n,u}$

Thank you!

- Mike Klaas, Nando De Freitas, and Arnaud Doucet. Toward practical N^2 Monte Carlo: The marginal particle filter. In *Proceedings of the 21st Conference in Uncertainty in Artificial Intelligence*, pages 308–315, 2005.
- Juan Kuntz, Francesca R Crucinio, and Adam M Johansen. Product-form estimators: exploiting independence to scale up Monte Carlo. *Statistics and Computing*, 32(1):1–22, 2022.
- Fredrik Lindsten, Adam M Johansen, Christian A Næsseth, Bonnie Kirkpatrick, Thomas B Schön, John A D Aston, and Alexandre Bouchard-Côté. Divide-and-Conquer with sequential Monte Carlo. *Journal of Computational and Graphical Statistics*, 26(2):445–458, 2017.
- Christian A Næsseth, Fredrik Lindsten, and Thomas B Schön. Nested sequential Monte Carlo methods. In *Proceedings of the 32nd International Conference on Machine Learning*, volume 37, pages 1292–1301. Proceedings of Machine Learning Research, 2015.