

# Divide-and-Conquer sequential Monte Carlo with applications to high dimensional filtering

---

Francesca R. Crucinio

SIAM-UQ 2024

Joint work with Adam M. Johansen

- 1** Motivation
- 2 Divide and Conquer SMC
- 3 Divide and Conquer SMC for Filtering

# Spatio-temporal Data





Figure 1: Drought Detection from Rain Precipitations (Næsseth et al., 2015).

- **Observe** precipitation (in millimeters) for each location and year
- **Recover** if there is a drought in that location


# Spatial State Space Model



- 1 Motivation
- 2 Divide and Conquer SMC**
- 3 Divide and Conquer SMC for Filtering

# Divide-and-Conquer SMC (DaC-SMC)

An extension of standard SMC in which the sequence of target distributions  $\{\hat{\gamma}_u\}_{u \in \mathbb{T}}$  evolves on a tree  $\mathbb{T}$  rather than on a line (Lindsten et al., 2017).



# Divide-and-Conquer SMC (DaC-SMC)

At each node  $u$  we have a set of  $N$  weighted particles  $\{X_u^i, W_u^i\}_{i=1}^N$  approximating  $\hat{\gamma}_u$ .

To evolve use standard SMC ingredients

- transition/mutation:  $X_u^i \sim M_u((X_{\ell(u)}^i, X_{r(u)}^i), \cdot)$
- reweighting:  $W_u^i \propto G_u(X_u^i)$
- resampling

with the addition of a **merging** step when two branches of the tree merge.

# The Merge Step

- **Aim:** Given two populations of weighted particles on the left and the right child,  $\{X_{\ell(u)}^i, W_{\ell(u)}^i\}_{i=1}^N$  and  $\{X_{r(u)}^i, W_{r(u)}^i\}_{i=1}^N$ , approximating  $\hat{\gamma}_{\ell(u)}$  and  $\hat{\gamma}_{r(u)}$  build an approximation of  $\hat{\gamma}_u$



# The Merge Step

- **Aim:** Given two populations of weighted particles on the left and the right child,  $\{X_{\ell(u)}^i, W_{\ell(u)}^i\}_{i=1}^N$  and  $\{X_{r(u)}^i, W_{r(u)}^i\}_{i=1}^N$ , approximating  $\hat{\gamma}_{\ell(u)}$  and  $\hat{\gamma}_{r(u)}$  build an approximation of  $\hat{\gamma}_u$

- **Ingredients:**

- ▶ the (weighted) product form estimator ([Kuntz et al., 2022](#))

$$\hat{\gamma}_{\ell(u)}^N \times \hat{\gamma}_{r(u)}^N = \frac{1}{N^2} \sum_{i_1=1}^N \sum_{i_2=1}^N W_{\ell(u)}^{i_1} W_{r(u)}^{i_2} \delta_{(X_{\ell(u)}^{i_1}, X_{r(u)}^{i_2})}$$


- ▶  $N^2$  mixture (importance) weights

$$m_u^{(i_1, i_2)} := W_{\ell(u)}^{i_1} W_{r(u)}^{i_2} \frac{\hat{\gamma}_u(X_{\ell(u)}^{i_1}, X_{r(u)}^{i_2})}{\hat{\gamma}_{\ell(u)}(X_{\ell(u)}^{i_1}) \hat{\gamma}_{r(u)}(X_{r(u)}^{i_2})}$$

- ▶ a resampling scheme to obtain  $\{\tilde{X}_u^i, 1/N\}_{i=1}^N$  from  $\{(X_{\ell(u)}^{i_1}, X_{r(u)}^{i_2}), m_u^{(i_1, i_2)}\}_{i_1, i_2=1}^N$

- 1 Motivation
- 2 Divide and Conquer SMC
- 3 Divide and Conquer SMC for Filtering**

# Divide and Conquer SMC for Filtering



**Figure 2:** Space decomposition for  $d = 8$ .

# Choice of Distributions

Fix  $n$ . We need  $\{\hat{\gamma}_{n,u}\}_{u \in \mathbb{T}}$ .

## Constraints:

- $\hat{\gamma}_{n,\mathfrak{R}}(x_n) \propto p(x_n | y_{1:n})$
- for  $u \neq \mathfrak{R}$ ,  $\hat{\gamma}_{n,u}$  only depends on a subset of  $x_n(1:d)$
- we can compute

$$\frac{\hat{\gamma}_{n,u}(x_{n,\ell(u)}, x_{n,r(u)})}{\hat{\gamma}_{n,\ell(u)}(x_{n,\ell(u)}) \hat{\gamma}_{n,r(u)}(x_{n,r(u)})}$$

# Choice of Distributions

Recall

$$p(x_n|y_{1:n}) \propto g_n(x_n, y_n) \int f_n(x_{n-1}, x_n) p(x_{n-1}|y_{1:n-1}) dx_{n-1}.$$

Use auxiliary functions  $f_{n,u}, g_{n,u}$  such that  $f_{n,\mathfrak{R}} = f_n$ ,  $g_{n,\mathfrak{R}} = g_n$  and for  $u \in \mathbb{T} \setminus \mathfrak{R}$ ,  $f_{n,u}$  and  $g_{n,u}$  serve as proxies for marginals of  $f_n, g_n$

$$\hat{\gamma}_{n,u}(x_{n,u}) = g_{n,u}(x_{n,u}, (y_n(i))_{i \in u}) \int f_{n,u}(x_{n-1}, x_{n,u}) \hat{\gamma}_{n-1,\mathfrak{R}}(x_{n-1}) dx_{n-1}$$

# Choice of Distributions

Recall

$$p(x_n | y_{1:n}) \propto g_n(x_n, y_n) \int f_n(x_{n-1}, x_n) p(x_{n-1} | y_{1:n-1}) dx_{n-1}.$$

Use auxiliary functions  $f_{n,u}, g_{n,u}$  such that  $f_{n,\mathfrak{R}} = f_n$ ,  $g_{n,\mathfrak{R}} = g_n$  and for  $u \in \mathbb{T} \setminus \mathfrak{R}$ ,  $f_{n,u}$  and  $g_{n,u}$  serve as proxies for marginals of  $f_n, g_n$

$$\hat{\gamma}_{n,u}(x_{n,u}) = g_{n,u}(x_{n,u}, (y_n(i))_{i \in u}) \int f_{n,u}(x_{n-1}, x_{n,u}) \hat{\gamma}_{n-1,\mathfrak{R}}(x_{n-1}) dx_{n-1}$$

The integral above is intractable, so we define approximate targets as in marginal particle filters (Klaas et al., 2005)

$$\hat{\gamma}_{n,u}(x_{n,u}) = g_{n,u}(x_{n,u}, (y_n(i))_{i \in u}) W_{n-1}^i \sum_{i=1}^N f_{n,u}(X_{n-1,\mathfrak{R}}^i, x_{n,u})$$

# Divide and Conquer SMC for Filtering

At time  $n$ :

- start with  $\{X_{n-1,\mathfrak{R}}^i, W_{n-1}^i\}_{i=1}^N$  approximating  $p(x_{n-1}|y_{1:n-1})$  at the leaf level
- define targets

$$\hat{\gamma}_{n,u}(x_{n,u}) = g_{n,u}(x_{n,u}, (y_n(i))_{i \in u}) W_{n-1}^i \sum_{i=1}^N f_{n,u}(X_{n-1,\mathfrak{R}}^i, x_{n,u})$$

- use DaC to move up the space from the leaves to the root, i.e. from  $d$  particle populations approximating 1-dimensional to one  $d$ -dimensional population  $\{X_{n,\mathfrak{R}}^i, W_n^i\}_{i=1}^N$  which approximates  $p(x_n|y_{1:n})$

# Linear Gaussian Model

$$\begin{aligned}f_1(x_1) &= \mathcal{N}(x_1; m_1, \Sigma_1) \\f_n(x_{n-1}, x_n) &= \mathcal{N}(x_n; Ax_{n-1}, \Sigma) \\g_n(x_n, y_n) &= \mathcal{N}(y_n; x_n, \sigma_y^2 Id_d),\end{aligned}$$


with  $\Sigma \in \mathbb{R}^{d \times d}$  a tridiagonal matrix.

We compare the RMSE

$$\text{RMSE}(x_n(i)) := \frac{\mathbb{E}[(\bar{x}_n(i) - \mu_{n,i})^2]}{\sigma_{n,i}^2}$$



# Linear Gaussian Model - RMSE



# Spatial Model

The components of  $X_t$  are indexed by the vertices  $v \in V$  of a lattice, where  $V = \{1, \dots, d\}^2$ , and


$$X_n(v) = X_{n-1}(v) + U_n(v), \quad U_n(v) \sim \mathcal{N}(0, \sigma_x^2)$$

and

$$g_t(x_n, y_n) \propto$$


$$\left[ 1 + \nu^{-1} \sum_{v \in V} \left( (y_n(v) - x_n(v)) \sum_{j: D(v,j) \leq r_y} \tau^{D(v,j)} (y_n(j) - x_n(j)) \right) \right]^{-(\nu + d^2)/2}.$$

## $2 \times 2$ grid




# $8 \times 8$ grid

(1, 1)



Runtime / s

(8, 6)



Runtime / s

 dac  dac-ada

## Pros

- No need for analytical form of  $f_n(x_n(i) \mid x_n(1 : i - 1))$
- No need for factorised likelihoods
- Easy to parallelise and distribute

## Cons

- Polynomial cost in  $N$  (can be mitigated via GPUs)
- Needs specification of  $\hat{\gamma}_{n,u}$

**Thank you!**

- Mike Klaas, Nando De Freitas, and Arnaud Doucet. Toward practical  $N^2$  Monte Carlo: The marginal particle filter. In *Proceedings of the 21st Conference in Uncertainty in Artificial Intelligence*, pages 308–315, 2005.
- Juan Kuntz, Francesca R Crucinio, and Adam M Johansen. Product-form estimators: exploiting independence to scale up Monte Carlo. *Statistics and Computing*, 32(1):1–22, 2022.
- Fredrik Lindsten, Adam M Johansen, Christian A Næsseth, Bonnie Kirkpatrick, Thomas B Schön, John A D Aston, and Alexandre Bouchard-Côté. Divide-and-Conquer with sequential Monte Carlo. *Journal of Computational and Graphical Statistics*, 26(2):445–458, 2017.
- Christian A Næsseth, Fredrik Lindsten, and Thomas B Schön. Nested sequential Monte Carlo methods. In *Proceedings of the 32nd International Conference on Machine Learning*, volume 37, pages 1292–1301. Proceedings of Machine Learning Research, 2015.